



Advanced topics in NoSQL  
databases  
ESILV - A4 IBO

# Introduction mongoDB

- Distributed NoSQL database
  - Large data management (Humongous)
  - Document-oriented NoSQL
    - JSON documents
    - (BSON Objects serialization)
- Key points:
  - Simple to integrate for developers
  - Rich query language
    - MQL: Mongo Query Language
  - Several optimization features
    - Attributes indexing (Shard/Btree/RTree)
  - Smart data allocation
    - Sharding (GridFS) [+ tagging / clustering]
  - Several deployment
    - Private/Public Cloud, Local (On Premise)

# Interactions: JavaScript

- JS objects
  - Attributes + functions
  - *db*: database
  - *sh*: sharding
  - *rs*: replica set
- MQL
  - "JSON" = pattern
- MapReduce
  - Functions (untyped)



Studio3t



Robo3t

# Starting Commands

- **Select database (shell)**

- Command: → **use myBD** ;

- **Collections** of documents

- Create: → **db.createCollection('users');**
- Manipulate: → **db.users.** <command> ;
  - Commands : **find(), save(), delete(), update(), aggregate(), distinct(), mapReduce()**...
  - Meaning in SQL: *FROM*

- **Documents**

- JSON:

```
{  
  "_id": ObjectId("4efa8d2b7d284dad101e4bc7"),  
  "name": "James Bond", "login": "james", "age": 50,  
  "address": {"street": "30 Wellington Square", "city": "London"},  
  "job" : ["agent", "MI6"]  
}
```

- Insert: → **db.users.save (    );**  **//no quotes!**

# MQL: find queries<sup>1</sup> (1/2)

- Document oriented queries
- Command: `→ db.users.find( <filter> , <projection> );`
- **Filter:**
  - “JSON” pattern that must be *matched* on the document
  - “Key/Value” format
  - Can embed exact matching, operations, nesting, arrays
    - Operations: `$op` (no quotes)
  - Meaning in SQL: *WHERE*
- **Projection:**
  - “Key/Value” pairs that are given in output
  - Meaning in SQL: *SELECT (no aggregates)*
- Example:  
`→ db.users.find( {“login” : “james”} , {“name” : 1, “age” : 1});`

<sup>1</sup> - find queries are simple queries for your report

# MQL: find queries (2/2)

- Exact match:

→ `db.users.find( { "login" : "james" }, { "name" : 1, "age" : 1 });`

- With nesting "." for the path to the given key

→ `db.users.find( { "address.city" : "London" } );`

- Operations

→ `db.users.find( { "age" : { $gt : 40 } } );`

`// $gt, $gte, $lt, $lte, $ne, $in, $nin, $or, $and, $exists, $type, $size, $cond...`

- Regular expressions<sup>1</sup>

→ `db.users.find( { "name" : { $regex : "james", $options : "i" } } );`

- Arrays

→ `db.users.find( { "job" : "MI6" } );`

`// Check in the list`

→ `db.users.find( { "job.1" : "MI6" } );`

`// 2nd position in the list`

→ `db.users.find( { "job" : ["MI6"] } );`

`// List exact match`

1 - regex : <https://docs.mongodb.com/manual/reference/operator/query/regex/>

# MQL: Distinct - Count

- List of distinct values from a key
  - `db.users.distinct( "name" );`
  - `db.users.distinct( "address.city" );`
- Number of documents
  - `db.users.count();`
  - `db.users.find( { "age" : 50 }).count();`

# MQL: pipeline<sup>2</sup> (1/3)

- **aggregate()** : Ordered sequence of operators *aggregation pipeline*
- Command:  
→ `db.users.aggregate( [ {$op1 : {}}, {$op2 : {}}, ... ] );`
- **Operators:**
  - `$match` : simple queries //meaning in SQL: WHERE
  - `$project` : projections //meaning in SQL: SELECT
  - `$sort` : sort result set //meaning in SQL: ORDER BY
  - `$unwind` : normalization in 1NF
  - `$group` : group by value + aggregate function // meaning in SQL: GROUP BY + fn
  - `$lookup` : left outer join (since 3.2 – work locally)
  - `$out` : store the output in a collection (since 3.2)
  - `$geoNear` : sort by nearest points (lat/long)
  - ...

2 – Pipeline queries (aggregate) are complex queries. For long pipelines, it can be hard queries



# MQL: pipeline (2/3)

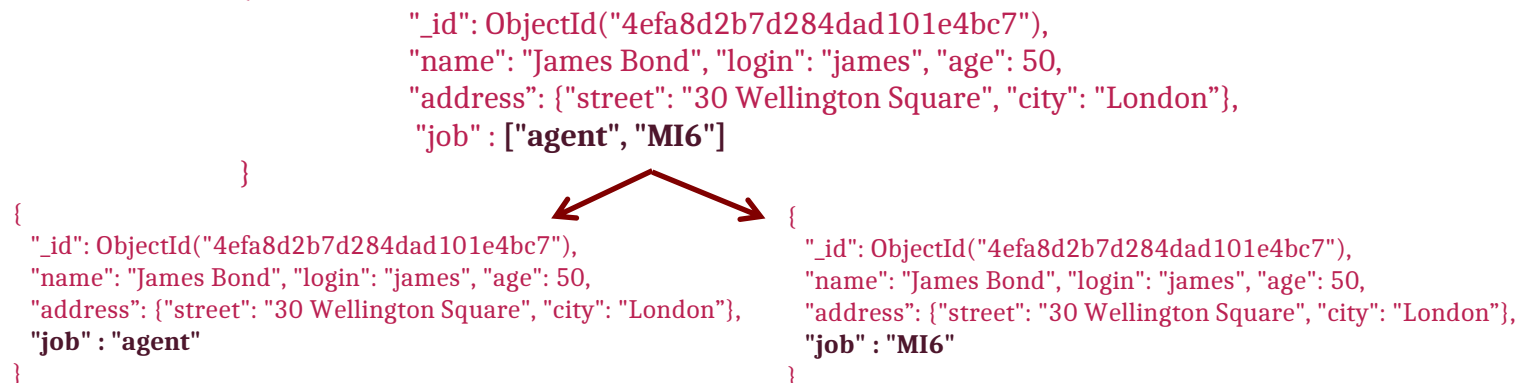
- **Pipeline** : Each step (operator) gives its output to the following operator (input)

```
> db.users.aggregate([{$match : {"address.city" : "London"}},
                      {$project : {"login" : 1, "age" : 1}},
                      {$sort : {"age" : 1, "login" : -1}}
                      ]);
```

- **\$unwind**

- Unnest an array and produce a new document for each item in the list

```
> db.users.aggregate([ {$unwind : "$job"} ]);
```



# MQL: pipeline (3/3)

- **\$group** : key (**\_id**) + aggregate (\$sum / \$avg / ...)  
No grouping key: *only one output*  
→ `db.users.aggregate([ {$group : {"_id" : "age", "res": {$sum : 1}}} ])`  
Group by value: **\$key**  
→ `db.users.aggregate([ {$group:{"_id" : "$age", "res": {$sum : 1}}} ])`  
Apply an aggregate function: **\$key**  
→ `db.users.aggregate([{$group:{"_id":"$address.city", "avg": {$avg: "$age"}}}] )`
- Sequence **example**  
→ `db.users.aggregate([  
 {$match: {"address.city" : "London"} },  
 {$unwind : "$job" },  
 {$group : {"_id" : "$job", "val": {$avg: "$age"}} },  
 {$match : {"val" : {$gt : 30}} },  
 {$sort : {"val" : -1} } ])`

# MQL: updates

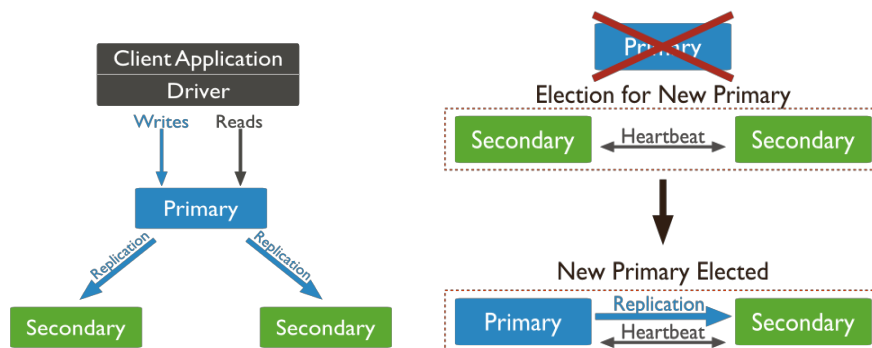
```
> db.users.update (  
    { "_id" : ObjectId("4efa8d2b7d284dad101e4bc7") },  
    { "$inc" : { "age" : 1 } } );
```

- Atomic updates (one document)
  - \$set – modify value
  - \$unset – remove a key/value
  - \$inc – Increment
  - \$push – Add inside an array
  - \$pushAll – Add several values
  - \$pull – Remove a value in an array
  - \$pullAll – Remove several values

*UpdateAll* – all documents are updated

# Replica Set

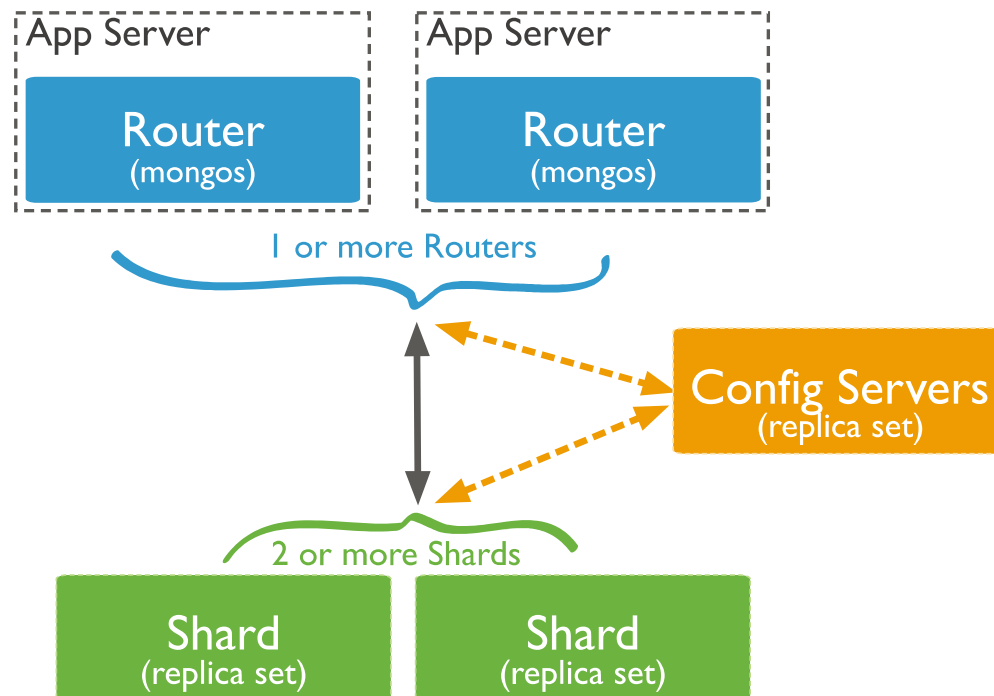
- Data replication
  - Asynchronous
    - Primary server: writes
    - Secondary servers: reads
  - Fault tolerance
    - New primary server election
    - Needs an arbiter
- Consistency vs availability
  - Reads applied on the primary (by default)
  - Updates via oPlog (log file)



# Sharding

- Data distribution on a cluster
- Data balancing
  - Sharding key choice: beginning
- 1 Shard = 1 Replica Set
- Min config:
  - *Config Servers (x3)*
  - *Mongos (router x3)*
- Partitioning key (sharding)
  - Ranged-based (GridFS : sort)
    - `sh.shardCollection( "myBD.users", "login" );`
  - Hash-based (md5 sur clé)
    - `sh.shardCollection( "myBD.users", { "_id": "hashed" } )`
  - By zones/tags (+ sharding)

# Sharding & Replica Sets



# Indexing

- Create a **BTree**

→ `db.users.createIndex( {"age":1} ) ;`

- All queries on « age » become more efficient
  - Even for Map/Reduce with « queryParam »
  - Execution plan « .explain() »
- No combination of indexes

# Indexing: 2DSphere

## Apply geolocalized queries

→ `db.users.ensureIndex( { "address.location" : "2dsphere" } );`

- Localization format

`"location":{"type": "Point", "coordinates" : [51.489220, -0.162866]}`

- Spherical queries

`var near = {$near: {$geometry:{"type":"Point","coordinates":[51.489220,-0.162866]},  
$maxDistance : 10000}};`

`db.users.find( {"address.location":near}, {"name":1, "_id":0});`

- \$geoNear operator (*aggregate*)

`var geoNear = {"near":{"type":"Point", "coordinates" : [51.489220,-0.162866]},  
"maxDistance":10000, "distanceField" : "outputDistance", "spherical":true};`

`db.users.aggregate([ {"$geoNear": geoNear} ]);`

- « polygonales » queries

`var polygon = {$geoWithin : {$geometry : {"type" : "Polygon", "coordinates" : [ [P1], [P2], [P3], [P1] ]}}}`

<http://docs.mongodb.org/manual/tutorial/query-a-2d-index/>




# Storage Engine

- **Wired Tiger**
  - By default since 3.2
  - Created for *BerkeleyDB*, used for Oracle NoSQL
  - Concurrency purpose: multi-version
- **MMAPV1**
  - Original engine
  - Mapping in main memory
  - Lock on collections (long writes)
- **Encrypted storage**
  - Since 3.2
  - Based on Wired Tiger
  - KMIP (Key Management Interoperability Protocol)
- **In Memory**
  - Since 3.2.6
  - Based on Wired Tiger
  - Must fit in main memory (error : *WT\_CACHE\_FULL*)

# Wired Tiger – Cache system

- 50% RAM – 1GB (or 256MB)
- Data are compressed in main memory



# Drivers / API

- Drivers: <http://docs.mongodb.org/ecosystem/drivers/>
  - Python, Ruby, Java, Javascript (Node.js), C++, C#, PHP, Perl, Scala...
  - Syntax: <http://docs.mongodb.org/ecosystem/drivers/syntax-table/>

# Driver Java

## Connection

```
ServerAddress server = new ServerAddress ("localhost",
27017);
```

```
Mongo mongo = new Mongo(server);
```

## Environment (DB + collection)

```
DB db = mongoClient.getDB( "maBD" );
```

```
DBCollection users = db.getCollection("users");
```

## Authentication

```
db.authenticate(login, passwd.toCharArray());
```

## Document inserting (DBObject)

```
DBObject doc = new BasicDBObject("name", "MongoDB")
```

```
.append("type", "database")
```

```
.append("count", 1)
```

```
.append("info", new BasicDBObject("x",
203).append("y", 102));
```

```
// or doc = (DBObject) JSON.parse(jsonText) ;
users.insert(doc);
```

## Querying (pattern query + cursor)

```
BasicDBObject query = new
```

```
BasicDBObject("name", "James Bond");
```

```
DBCursor cursor = coll.find(query);
```

```
try {
```

```
while(cursor.hasNext()) {
```

```
System.out.println(cursor.next());
```

```
}
```

```
} finally {
```

```
cursor.close();
```

```
}
```

## Map/Reduce :

```
MapReduceCommand cmd =
```

```
new MapReduceCommand("users", map,
```

```
reduce, "outputColl",
```

```
MapReduceCommand.
```

```
OutputType.REPLACE, query);
```

```
MapReduceOutput out = users.mapReduce(cmd);
```

```
for (DBObject o : out.results()) {
```

```
System.out.println(o.toString());
```

```
}
```

```
//outputType : INLINE, REPLACE, MERGE, REDUCE
```

[http://api.mongodb.org/java/current/index.html?\\_ga=1.177444515.761372013.1398850293](http://api.mongodb.org/java/current/index.html?_ga=1.177444515.761372013.1398850293)

# Driver C#

## References/Libraries

```
MongoDB.Bson.dll
    using MongoDB.Bson;
MongoDB.Driver.dll
    using MongoDB.Driver;
```

## Connection

```
var connectionString = "mongodb://localhost";
var client = new MongoClient(connectionString);
var server = client.GetServer();
```

## Database & Collection

```
var db = server.GetDatabase("maBD");
var coll = db.GetCollection<User>("users");
```

## Objet « User » to define

```
public class User{
    public ObjectId Id { get; set; }
    public string name { get; set; }
    public string login { get; set; }
    public int age { get; set; }
    public Address address { get; set; }
}
```

## Get a document :

```
var query=Query<User>.EQ(e=>e.login,"james");
var entity = coll.FindOne(query);
```

## Package for output results: LINQ

```
using MongoDB.Driver.Linq;
```

## Queries:

```
var query = coll.AsQueryable<User>()
    .Where(e => e.age > 40)
    .OrderBy(c => c.name);
```

## Answer:

```
foreach (var user in query)
{
    // traitement
}
```

## MapReduce :

```
var mr = coll.MapReduce(map, reduce);
foreach (var document in mr.GetResults()) {
    Console.WriteLine(document.ToJson());
}
```

<http://www.nuget.org/packages/mongocsharpdriver/>

# Driver Python

## Importation

```
>>> import pymongo
```

## Connexion

```
>>> from pymongo import MongoClient
>>> client = MongoClient()
>>> client = MongoClient('localhost', 27017)
```

## Database & collection

```
>>> db = client.maBD
>>> coll = db.users
```

## GET one document

```
coll.find_one ( { "login" : "james" } )
```

## Query

```
>>> for c in coll.find ( { "age": 40 } ) :
...     pprint.pprint (c)
```

## MapReduce

```
>>> from bson.code import Code
>>> map = Code("function () {"
... "  this.tags.forEach(function(z) {"
... "    emit(z, 1);"
... "  });"
... "}")
>>> reduce = Code("function (key, values) {"
... "  var total = 0;"
... "  for (var i = 0; i < values.length; i++) {"
... "    total += values[i];}"
... "  return total; } ")
>>> result = coll.map_reduce(map, reduce,
...                           "myresults")
>>> for doc in myresults.find():
...     print doc
```